

Single parameter models

Dhafer Malouche

Probability distributions

- Bernoulli Distribution

- Binomial Distribution

- Poisson Distribution

- Negative Binomial Distribution

- Normal Distribution

- Exponential Distribution

- Gamma Distribution

- Beta Distribution

Bayesian single parameter Models

- Bernoulli Model

- Normal Model

- Poisson Model

Probability distributions

Bernoulli Distribution

Bernoulli Distribution: Definition

- ▶ Discrete probability distribution for **binary outcomes**
- ▶ Random variable X takes values:

$$X = \begin{cases} 1 & \text{with probability } p \\ 0 & \text{with probability } 1 - p \end{cases}$$

- ▶ Probability mass function (PMF):

$$P(X = k) = p^k(1 - p)^{1-k} \quad \text{for } k \in \{0, 1\}$$

- ▶ Parameters: $p \in [0, 1]$ (probability of success)

Real-World Example

Login Attempt Success

Consider a website login system where:

- ▶ Each login attempt succeeds with probability $p = 0.3$
- ▶ Fails with probability 0.7

Let:

$$X = \begin{cases} 1 & \text{Successful login} \\ 0 & \text{Failed login} \end{cases}$$

This is a Bernoulli($p = 0.3$) random variable.

Simulation and Comparison in R

```
# Parameters
n <- 1000
p <- 0.3

# Generate random samples
set.seed(123)
samples <- rbinom(n, 1, p)

# Calculate empirical probabilities
emp_probs <- table(samples)/n
```

Theoretical results:

- ▶ $P(X=0) = 0.7$
- ▶ $P(X=1) = 0.3$

Empirical results:

- ▶ $P(X=0) = \text{emp_probs}[1]$
- ▶ $P(X=1) = \text{emp_probs}[2]$

PMF Comparison:

```
# Theoretical PMF
```

```
theoretical <- c(0.7, 0.3)
```

```
# Plot comparison
```

```
barplot(rbind(theoretical, emp_probs),  
        beside = TRUE,  
        col = c("red", "blue"),  
        main = "PMF Comparison")
```


CDF Comparison:

```
# Plot theoretical CDF
plot(stepfun(c(0,1), c(0,0.7,1)),
     main = "CDF Comparison",
     xlim = c(-0.5,1.5))

# Add empirical CDF
lines(ecdf(samples),
     col = 'blue', lty = 2)
```

Activity: Simulating Disease Transmission

Scenario: COVID-19 exposure at a school ($N = 500$ students)

Simulation Steps:

```
# Set parameters:
set.seed(456)
# Transmission probability
p_infect <- 0.15
n_students <- 500
# Simulate infections:
infected <- rbinom(n_students, 1, p_infect)
# Calculate outcomes:
n_infected <- sum(infected)
attack_rate <- mean(infected)
# Visualize
barplot(c(Theoretical = p_infect,
          Observed = attack_rate),
       ylab = "Probability",
       col = c("red3", "dodgerblue"),
       main = "Infection Risk Comparison")
```

Immediate Tasks:

- ▶ Report number infected
- ▶ Compare red vs. blue bars
- ▶ Compute: $\frac{\text{Infected}}{\text{Total}}$
- ▶ Expected vs. observed cases

Activity: Analysis & Public Health Insights (Part 1)

Follow-up from Simulation Results

```
# Confidence interval for attack rate  
prop.test(n_infected , n_students )
```

```
# Simulate new variant impact  
p_new_variant <- 0.30 # 30% infection probability  
infected_new <- rbinom(n_students , 1, p_new_variant )
```

Initial Questions:

- ▶ How would you interpret the 95% confidence interval?
- ▶ What does changing p_new_variant from 0.15 to 0.30 imply about:
 - ▶ Hospitalization rates?
 - ▶ Healthcare system capacity?

Key concepts introduced: Confidence intervals, parameter sensitivity

Activity: Analysis & Public Health Insights (Part 2)

Vaccine Efficacy (VE)

- ▶ Definition: Relative reduction in infection risk for vaccinated vs unvaccinated
- ▶ Formula:

$$VE = 1 - \frac{\text{Risk}_{\text{vaccinated}}}{\text{Risk}_{\text{unvaccinated}}}$$

- ▶ R implementation:

```
p_control <- 0.25 # Infection rate in control group
p_vax <- 0.05 # Infection rate in vaccinated group
ve <- 1 - (p_vax/p_control) # Returns 0.8 (80% efficacy)
```

Sample Size Calculation

- ▶ Theory: For 95% CI with margin of error e: We would like to know the needed sample if we would a 95% CI mean $\pm 2\%$

$$n = \frac{Z_{\alpha/2}^2 \cdot p(1-p)}{e^2}$$

- ▶ R implementation:

```
p <- 0.5 # Conservative estimate
e <- 0.02 # Desired precision
n <- (qnorm(0.975)^2 * p*(1-p)) / e^2
```

Binomial Distribution

Binomial Distribution: Definition

- ▶ Discrete distribution for **success counts in n trials**
- ▶ $X \sim \text{Bin}(n, p)$ where:
 - ▶ n : Number of independent trials
 - ▶ p : Success probability per trial
- ▶ PMF:

$$P(X = k) = \binom{n}{k} p^k (1 - p)^{n-k} \quad k = 0, 1, \dots, n$$

- ▶ Expectation: $\mathbb{E}[X] = np$
Variance: $\text{Var}(X) = np(1 - p)$

Vaccine Efficacy Trial

In a COVID-19 vaccine trial with 100 participants:

- ▶ Each participant has 5% infection risk (placebo group)
- ▶ Let X = number of infections in the group
- ▶ $X \sim \text{Bin}(n = 100, p = 0.05)$

Probability questions:

- ▶ $P(X \geq 10)$ - Extreme outbreak risk
- ▶ $P(5 \leq X \leq 15)$ - Expected range

Simulation and Comparison in R

```
# Parameters
n_trials <- 1000
n <- 100    # Participants
p <- 0.05   # Infection risk

# Simulate outbreaks
set.seed(123)
infections <- rbinom(n_trials, n, p)

# Theoretical PMF
k <- 0:n
pmf <- dbinom(k, n, p)
```

Comparison checks:

- ▶ Empirical mean vs np
- ▶ Sample variance vs $np(1 - p)$
- ▶ Histogram shape vs PMF

```
mean(infections)
# Should be ~5
var(infections)
# Should be ~4.75
```


► PMF Comparison:

```
hist(infections, freq=FALSE,  
     main = "Infection Distribution",  
     xlab = "Number of Cases")  
lines(k, pmf, col="red", lwd=2)
```

► CDF Comparison:

```
plot(ecdf(infections),  
     main = "CDF Comparison")  
lines(k, pbinom(k, n, p),  
      col="red", type="s")
```

Poisson Distribution

Poisson Distribution: Definition

- ▶ Discrete distribution for **event counts in fixed interval/area**
- ▶ $X \sim \text{Pois}(\lambda)$ where:
 - ▶ λ : Average rate (events per unit)
 - ▶ Constant event risk, independent occurrences

- ▶ PMF:

$$P(X = k) = \frac{e^{-\lambda} \lambda^k}{k!} \quad k = 0, 1, 2, \dots$$

- ▶ Expectation/Variance: $\mathbb{E}[X] = \text{Var}(X) = \lambda$

Microplastic Pollution Monitoring

Water quality researchers study a river:

- ▶ Average of 4 microplastic particles per liter ($\lambda = 4$)
- ▶ X = number of particles in 1L sample
- ▶ $X \sim \text{Pois}(4)$

Key questions:

- ▶ $P(X \geq 8)$ - Extreme pollution probability
- ▶ $P(X \leq 2)$ - Compliance with safety standards
- ▶ 95% prediction interval for particle counts

Simulation and Comparison in R

```
# Parameters
lambda <- 4
n_samples <- 1000

# Simulate water samples
set.seed(123)
particles <- rpois(n_samples, lambda)

# Theoretical PMF
k <- 0:15
pmf <- dpois(k, lambda)
```

Validation checks:

- ▶ Empirical mean:
mean(particles)
- ▶ Sample variance:
var(particles)
- ▶ Zero-inflation test:
mean(particles==0)

```
# Expected
# vs observed zeros
exp(-lambda)
# Theoretical P(X=0)
mean(particles == 0)
```

► PMF Comparison:

```
hist(particles, freq=FALSE, breaks=0:15,  
     main = "Microplastic Particles/Liter",  
     xlab = "Count")  
points(k, pmf, col="red", pch=19)
```

► Time Series:

```
plot(particles[1:50], type="b",  
     main = "Particle Count Time Series",  
     xlab = "Sample ID", ylab = "Count",  
     col = "darkgreen")  
abline(h = lambda, col="red", lty=2)
```

Activity: Pollution Hotspot Analysis

Scenario: Comparing two river sites

Simulation Task:

```
# Downstream (polluted)
lambda_ds <- 8
# Upstream (reference)
lambda_us <- 3

samples <- 500
set.seed(456)
ds_counts <- rpois(samples, lambda_ds)
us_counts <- rpois(samples, lambda_us)
```

Analysis:

- ▶ Test $\lambda_{ds} > \lambda_{us}$
- ▶ Compute $P(ds \geq 10)$
- ▶ Visualize PMF comparisons
- ▶ Estimate pollution ratio
- ▶ Check dispersion (var/mean ratio)

Negative Binomial Distribution

Negative Binomial Distribution: Definition (Traditional)

- ▶ Discrete distribution for the number of trials needed to achieve r successes.
- ▶ Assumes independent Bernoulli trials with success probability p .
- ▶ PMF:

$$P(X = k) = \binom{k-1}{r-1} p^r (1-p)^{k-r}, \quad k = r, r+1, r+2, \dots$$

- ▶ Moments:
 - ▶ Mean: $\mathbb{E}[X] = \frac{r}{p}$
 - ▶ Variance: $\text{Var}(X) = \frac{r(1-p)}{p^2}$

Customer Acquisition Costs

A SaaS company models trial-to-paid conversions:

- ▶ Each website visit is a trial with conversion probability $p = \frac{1}{3}$.
- ▶ A conversion requires one successful visit ($r = 1$).
- ▶ X = number of visits until conversion.
- ▶ $X \sim \text{NB}(r = 1, p = \frac{1}{3})$ (geometric distribution).

Key business questions:

- ▶ What is the probability of conversion in ≤ 2 visits?
- ▶ What is the 95th percentile of required visits?
- ▶ How does this inform customer acquisition cost modeling?

Simulation and Comparison in R

```
# Parameters
r <- 1          # Conversion requires 1 success
p <- 1/3        # Conversion probability per visit
n <- 1000

# Simulate number of visits until conversion
set.seed(2023)
failures <- rnbinom(n, size = r, prob = p)
visits <- failures + r # Total visits (failures + 1 success)

# Theoretical PMF for number of visits
k <- 1:15
pmf <- p * (1 - p)^(k - 1)

# Empirical dispersion check
dispersion <- var(visits) / mean(visits)
dispersion # Should be > 1 indicating overdispersion
```

- ▶ PMF Comparison:

```
hist(visits, freq=FALSE, breaks=0:20,  
      main = "Visits Until Conversion",  
      xlab = "Website Visits")  
lines(k, pmf, col="red", type="h")
```

- ▶ QQ-Plot:

```
# Compare to Poisson  
qqplot(rpois(1000, mu), visits,  
        main = "NB vs Poisson QQ-Plot")  
abline(0,1, col="red")
```

Normal Distribution

Normal Distribution: Definition

- ▶ Continuous distribution for **symmetric, bell-shaped data**
- ▶ $X \sim N(\mu, \sigma^2)$ where:
 - ▶ μ : Mean (location parameter)
 - ▶ σ : Standard deviation (scale parameter)
- ▶ Probability density function (PDF):

$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2}$$

- ▶ Properties:
 - ▶ Symmetric about μ
 - ▶ 68-95-99.7 rule for standard normal

IQ Score Distribution

Standardized intelligence testing:

- ▶ IQ scores follow $N(100, 15^2)$
- ▶ Population parameters:
 - ▶ Mean $\mu = 100$
 - ▶ SD $\sigma = 15$
- ▶ Diagnostic thresholds:
 - ▶ Gifted: > 130 (top 2.3%)
 - ▶ Intellectual disability: < 70 (bottom 2.3%)

Simulation and Comparison in R

```
# Parameters
mu <- 100
sigma <- 15
n <- 1000

# Generate scores
set.seed(123)
iq_scores <- rnorm(n, mu, sigma)

# Theoretical PDF
x <- seq(55, 145, length=100)
pdf <- dnorm(x, mu, sigma)
```

Validation checks:

- ▶ Empirical mean/SD vs parameters
- ▶ Shapiro-Wilk normality test
- ▶ Check 68-95-99.7 rule

```
mean(iq_scores)
# Should be ~100
sd(iq_scores)
# Should be ~15
shapiro.test(iq_scores)
```


- Density Comparison:

```
hist(iq_scores, freq=FALSE, breaks=20,  
     main = "IQ Score Distribution",  
     xlab = "IQ")  
lines(x, pdf, col="red", lwd=2)
```

- QQ-Plot:

```
qqnorm(iq_scores, main = "Normal Q-Q Plot")  
qqline(iq_scores, col="red")
```

Exponential Distribution

Exponential Distribution: Definition

- ▶ Continuous distribution for **time between events in Poisson process**
- ▶ $X \sim \text{Exp}(\lambda)$ where:
 - ▶ λ : Rate parameter (events per unit time)
 - ▶ Mean time between events: $\mu = 1/\lambda$
- ▶ PDF:

$$f(x) = \lambda e^{-\lambda x} \quad x \geq 0$$

- ▶ Key property: Memoryless

$$P(X > s + t | X > s) = P(X > t)$$

Network Intrusion Detection

A corporate network experiences:

- ▶ Average of 1 intrusion every 10 days ($\mu = 10$)
- ▶ $\lambda = 0.1$ intrusions/day
- ▶ X = days between intrusions
- ▶ $X \sim \text{Exp}(0.1)$

Security questions:

- ▶ $P(\text{Next intrusion} < 24\text{hrs}) = ?$
- ▶ 90th percentile of safe period
- ▶ Probability of > 30 days without intrusion

Simulation and Comparison in R

```
# Parameters
lambda <- 0.1 # Rate
mu <- 1/lambda # Mean = 10
n <- 1000

# Simulate inter-arrival times
set.seed(123)
durations <- rexp(n, lambda)

# Theoretical PDF
x <- seq(0, 30, length=100)
pdf <- dexp(x, lambda)
```

Validation checks:

- ▶ Empirical mean vs $1/\lambda$
- ▶ Variance vs $1/\lambda^2$
- ▶ Memoryless property test

```
mean(durations)
# Should be ~10
var(durations)
# Should be ~100
```

► Density Comparison:

```
hist(durations, freq=FALSE, breaks=30,  
     main = "Time Between Intrusions",  
     xlab = "Days")  
lines(x, pdf, col="red", lwd=2)
```

► Survival Function:

```
plot(ecdf(durations),  
     main = "CDF Comparison")  
lines(x, pexp(x, lambda),  
     col="red", lwd=2)
```

Gamma Distribution

Gamma Distribution: Definition

- ▶ Continuous distribution for **positive-valued, skewed data**
- ▶ $X \sim \Gamma(k, \theta)$ where:
 - ▶ k : Shape parameter (controls skewness)
 - ▶ θ : Scale parameter
- ▶ PDF:

$$f(x) = \frac{1}{\Gamma(k)\theta^k} x^{k-1} e^{-x/\theta} \quad x > 0$$

- ▶ Properties:
 - ▶ $\mathbb{E}[X] = k\theta$, $\text{Var}(X) = k\theta^2$
 - ▶ Generalizes Exponential ($k = 1$) and χ^2 distributions

Healthcare Example

Hospital Length of Stay

Modeling patient stays in a surgical ward:

- ▶ Average stay duration: 5 days ($k = 2, \theta = 2.5$)
- ▶ X = number of days hospitalized
- ▶ $X \sim \Gamma(2, 2.5)$

Clinical applications:

- ▶ $P(\text{Stay} > 10 \text{ days})$ - Identify long-stay patients
- ▶ 75th percentile for resource planning
- ▶ Compare recovery times between procedures
- ▶ Model healthcare costs associated with stays

Why Gamma?

Simulation and Comparison in R

Parameters

```
shape <- 2  
scale <- 2.5  
n <- 1000
```

Simulate hospital stays

```
set.seed(123)  
stays <- rgamma(n, shape=shape, scale=scale)
```

Theoretical PDF

```
x <- seq(0, 20, length=100)  
pdf <- dgamma(x, shape=shape, scale=scale)
```

Validation checks:

- ▶ Empirical mean vs $k\theta$
- ▶ Variance vs $k\theta^2$
- ▶ Skewness assessment

```
mean(stays)  
# Should be ~5  
var(stays)  
# Should be ~12.5  
moments::skewness(stays)  
# ~1.41
```

► Density Comparison:

```
hist(stays, freq=FALSE, breaks=30,  
     main = "Hospital Stay Duration",  
     xlab = "Days", col="lightblue")  
lines(x, pdf, col="maroon", lwd=2)  
legend("topright",  
       legend=c("Theory", "Empirical"),  
       col=c("maroon", "lightblue"), lwd=2)
```

► CDF Comparison:

```
plot(ecdf(stays),  
     main="Empirical vs Theoretical CDF")  
lines(x, pgamma(x, shape=shape, scale=scale),  
      col="darkgreen", lwd=2)
```

Beta Distribution

Beta Distribution: Definition

- ▶ Continuous distribution for **probabilities/proportions** (0-1)
- ▶ $X \sim \text{Beta}(\alpha, \beta)$ where:
 - ▶ α : Shape 1 (successes + 1)
 - ▶ β : Shape 2 (failures + 1)

- ▶ PDF:

$$f(x) = \frac{x^{\alpha-1}(1-x)^{\beta-1}}{B(\alpha, \beta)} \quad 0 \leq x \leq 1$$

- ▶ Properties:

- ▶ $\mathbb{E}[X] = \frac{\alpha}{\alpha + \beta}$
- ▶ $\text{Var}(X) = \frac{\alpha\beta}{(\alpha + \beta)^2(\alpha + \beta + 1)}$
- ▶ Generalizes Uniform ($\alpha = \beta = 1$)

A/B Test Conversion Rates

Comparing website versions:

- ▶ Version A: 45 conversions / 1000 visits
- ▶ Version B: 65 conversions / 1000 visits
- ▶ Model conversion rates as:
 - ▶ $r_A \sim \text{Beta}(46, 956)$
 - ▶ $r_B \sim \text{Beta}(66, 936)$

Why Beta?

- ▶ Natural for bounded probabilities
- ▶ Conjugate prior for Binomial
- ▶ Flexible uncertainty representation

Business questions:

- ▶ $P(r_B > r_A)$ - Version superiority probability
- ▶ 95% credible intervals for rates
- ▶ Minimum detectable effect size

Simulation and Comparison in R

```
# Parameters
alpha <- 46
beta <- 956
n <- 10000

# Simulate conversion rates
set.seed(123)
rates <- rbeta(n, alpha, beta)

# Theoretical PDF
x <- seq(0, 0.1, length=100)
pdf <- dbeta(x, alpha, beta)
```

Validation checks:

- ▶ Empirical mean vs $\alpha/(\alpha + \beta)$
- ▶ Variance comparison
- ▶ Uniform check when $\alpha = \beta = 1$

```
mean(rates)
# Should be ~0.046
quantile(rates, c(0.025, 0.975))
```

► Density Comparison:

```
hist(rates, freq=FALSE, breaks=50,  
     main = "Conversion Rate Distribution",  
     xlab = "Conversion Rate", col="skyblue")  
lines(x, pdf, col="darkred", lwd=2)
```


► CDF Comparison:

```
plot(ecdf(rates), main="Beta CDF Comparison")  
lines(x, pbeta(x, alpha, beta),  
      col="darkgreen", lwd=2)
```


Bayesian single parameter Models

Single Parameter Mode

Bayesian Inference Process



$$\underbrace{p(\theta \mid \text{Data})}_{\text{posterior}} \propto \underbrace{L(\text{Data} \mid \theta)}_{\text{likelihood}} \times \underbrace{p(\theta)}_{\text{prior}}$$

Bernoulli Model

Example: Influence of Prior Beliefs on Bayesian Updating

- ▶ Scenario: Estimating the probability of success for a new treatment based on observed trial results.
- ▶ Observations: Suppose you observe 10 trials of a new treatment, where 6 trials are successful.
- ▶ Prior Beliefs:
 - ▶ **Non-Informative Prior:** We start with no prior information, reflecting a neutral stance on the treatment's effectiveness. This prior does not influence the likelihood gained from the observed data.
 - ▶ **Weakly Informative Prior:** We have some previous experience or expert opinion suggesting the treatment's success rate is around 50

Choosing the Prior Distribution

- ▶ **Prior:** Represents initial beliefs about p (probability of success).
- ▶ Common choices:
 - ▶ **Conjugate Prior (Beta):**

$$p \sim \text{Beta}(\alpha, \beta)$$

- α : Prior successes + 1
- β : Prior failures + 1
- ▶ **Non-Informative Prior:** $\text{Beta}(1, 1)$ (uniform over $[0, 1]$)
- ▶ **Weakly Informative Prior:** $\text{Beta}(2, 2)$ (gentle nudge toward 0.5)

Given observed data $y = (y_1, y_2, \dots, y_n)$, the posterior is:

$$p \mid y \sim \text{Beta}(\alpha + \text{successes}, \beta + \text{failures})$$

- ▶ **Successes:** $\sum_{i=1}^n y_i$
- ▶ **Failures:** $n - \sum_{i=1}^n y_i$

R Example: Analytical Solution

Data

```
successes <- 7  
failures <- 3
```

Priors

```
prior_uniform <- c(1, 1)           # Non-informative  
prior_weak <- c(2, 2)               # Weakly informative  
prior_informative <- c(5, 5)       # Informative
```

Posteriors

```
posterior_uniform <- c(prior_uniform[1] + successes,  
                        prior_uniform[2] + failures)  
posterior_weak <- c(prior_weak[1] + successes,  
                     prior_weak[2] + failures)  
posterior_informative <- c(prior_informative[1] + successes,  
                             prior_informative[2] + failures)
```

R Example: Analytical Solution

```
# Plot
curve(dbeta(x, posterior_uniform[1], posterior_uniform[2]),
      xlab="p", ylab="Density", col="blue",
      lwd=2, ylim=c(0, 4))
curve(dbeta(x, posterior_weak[1],
            posterior_weak[2]),
      col="red", lwd=2, add=TRUE)
curve(dbeta(x, posterior_informative[1],
            posterior_informative[2]),
      col="green", lwd=2, add=TRUE)
legend("topright",
      legend=c("Uniform", "Weak", "Informative"),
      col=c("blue", "red", "green"), lwd=2)
```


R Example: Using rstan

Stan Code (bernoulli_model.stan):

```
data {  
  int<lower=0> N;           // Number of trials  
  int<lower=0, upper=1> y[N]; // Binary outcomes  
}  
parameters {  
  real<lower=0, upper=1> p; // Probability of success  
}  
model {  
  p ~ beta(1, 1);           // Uniform prior  
  y ~ bernoulli(p);         // Likelihood  
}
```

R Example: Using rstan

R Code to Run Stan:

```
library(rstan)
y <- c(1, 0, 1, 1, 0, 1, 1, 1, 0, 1)
stan_data <- list(N = length(y), y = y)
fit <- stan(file = "bernoulli_model.stan",
            data = stan_data,
            iter = 2000, chains = 4)
print(fit)    # Summary of posterior
plot(fit)     # Visualize posterior distribution
```

Activity: Vaccine Efficacy

Scenario: A new vaccine is tested on 100 individuals. 15 develop side effects.

Task:

- ▶ Define a Beta prior for side effect probability p .
- ▶ Compute the posterior distribution.
- ▶ Compare results for:
 - ▶ Non-informative prior (Beta(1, 1))
 - ▶ Weakly informative prior (Beta(2, 2))
 - ▶ Informative prior (Beta(5, 95))
- ▶ Use `rstan` to fit the model and compare results.

Questions:

- ▶ What is $P(p > 0.2)$ under each prior?
- ▶ How does the choice of prior affect the posterior?
- ▶ What sample size would reduce prior influence?

Normal Model

Case 1: Unknown Mean, Known Variance

- ▶ **Model:** $y \sim N(\mu, \sigma^2)$, where σ^2 is known.
- ▶ **Conjugate Prior:** Normal distribution for μ :

$$\mu \sim N(\mu_0, \tau_0^2)$$

- μ_0 : Prior mean
- τ_0^2 : Prior variance

- ▶ **Posterior:**

$$\mu \mid y \sim N(\mu_n, \tau_n^2)$$

where:

$$\mu_n = \frac{\frac{\mu_0}{\tau_0^2} + \frac{n\bar{y}}{\sigma^2}}{\frac{1}{\tau_0^2} + \frac{n}{\sigma^2}}, \quad \tau_n^2 = \left(\frac{1}{\tau_0^2} + \frac{n}{\sigma^2} \right)^{-1}$$

R Example: Unknown Mean, Known Variance

```
# Data
y <- c(5.1, 5.5, 4.9, 5.3, 5.7) # Sample data
n <- length(y)
sigma2 <- 0.5 # Known variance

# Prior
mu0 <- 5.0 # Prior mean
tau02 <- 1.0 # Prior variance

# Posterior
mu_n <- (mu0/tau02 + n*mean(y)/sigma2) / (1/tau02 + n/sigma2)
tau_n2 <- 1 / (1/tau02 + n/sigma2)

# Plot
curve(dnorm(x, mu0, sqrt(tau02)), xlim=c(4, 6), ylab="Density",
      col="blue", lwd=2, main="Prior vs Posterior")
curve(dnorm(x, mu_n, sqrt(tau_n2)), col="red", lwd=2, add=TRUE)
legend("topright", legend=c("Prior", "Posterior"),
      col=c("blue", "red"), lwd=2)
```

Case 2: Unknown Mean and Variance

- ▶ **Model:** $y \sim N(\mu, \sigma^2)$, both μ and σ^2 unknown.
- ▶ **Conjugate Prior:** Normal-Inverse-Gamma (NIG):

$$\mu \mid \sigma^2 \sim N(\mu_0, \sigma^2 / \kappa_0), \quad \sigma^2 \sim \text{Inv-Gamma}(\alpha_0, \beta_0)$$

- μ_0 : Prior mean
- κ_0 : Prior precision scaling
- α_0, β_0 : Shape and scale for σ^2

- ▶ **Posterior:**

$$\mu \mid \sigma^2, y \sim N(\mu_n, \sigma^2 / \kappa_n), \quad \sigma^2 \mid y \sim \text{Inv-Gamma}(\alpha_n, \beta_n)$$

where:

$$\mu_n = \frac{\kappa_0 \mu_0 + n \bar{y}}{\kappa_0 + n}, \quad \kappa_n = \kappa_0 + n, \quad \alpha_n = \alpha_0 + \frac{n}{2},$$

$$\beta_n = \beta_0 + \frac{1}{2} \sum_{i=1}^n (y_i - \bar{y})^2 + \frac{\kappa_0 n (\bar{y} - \mu_0)^2}{2(\kappa_0 + n)}$$

R Example: Unknown Mean and Variance

```
# Data
y <- c(5.1, 5.5, 4.9, 5.3, 5.7) # Sample data
n <- length(y)

# Prior
mu0 <- 5.0 # Prior mean
kappa0 <- 1 # Prior precision scaling
alpha0 <- 2 # Prior shape for sigma^2
beta0 <- 1 # Prior scale for sigma^2

# Posterior
mu_n <- (kappa0 * mu0 + n * mean(y)) / (kappa0 + n)
kappa_n <- kappa0 + n
alpha_n <- alpha0 + n/2
beta_n <- beta0 + 0.5*sum((y - mean(y))^2) +
  (kappa0 * n*(mean(y) - mu0)^2)/(2*(kappa0 + n))
```


R Example: Unknown Mean and Variance

```
# Plot posterior for mu
curve(dnorm(x, mu_n, sqrt(beta_n / (alpha_n * kappa_n))),
      xlim=c(4, 6), ylab="Density", col="red", lwd=2,
      main="Posterior for Mean")
```

Using rstan for Unknown Mean and Variance

Stan Code (normal_model.stan):

```
data {  
  int<lower=0> N;          // Number of observations  
  real y[N];              // Data  
}  
parameters {  
  real mu;                // Mean  
  real<lower=0> sigma2;    // Variance  
}  
model {  
  mu ~ normal(5.0, sqrt(sigma2)); // Prior for mu  
  sigma2 ~ inv_gamma(2, 1);       // Prior for sigma2  
  y ~ normal(mu, sqrt(sigma2));   // Likelihood  
}
```

Using rstan for Unknown Mean and Variance

R Code to Run Stan:

```
library(rstan)
y <- c(5.1, 5.5, 4.9, 5.3, 5.7) # Sample data
stan_data <- list(N = length(y), y = y)
fit <- stan(file = "normal_model.stan", data = stan_data, ite
print(fit) # Summary of posterior
plot(fit) # Visualize posterior distributions
```

Activity: Climate Change Analysis

Scenario: Analyze annual temperature anomalies (in °C) for a city:

$$y = \{0.9, 1.1, 1.3, 1.0, 1.2, 1.4, 1.1, 1.3, 1.5, 1.2\}$$

Task:

- ▶ Assume $\sigma^2 = 0.1$ is known. Use a $\text{Normal}(1.0, 0.5)$ prior for μ . Compute the posterior.
- ▶ Assume both μ and σ^2 are unknown. Use a $\text{NIG}(1.0, 1, 2, 1)$ prior. Compute the posterior.
- ▶ Compare results using `rstan`.

Questions:

- ▶ What is $P(\mu > 1.2)$ under each model?
- ▶ How does prior choice affect results?
- ▶ What sample size reduces prior influence?

Poisson Model

Introduction to Poisson Model

- ▶ The Poisson distribution is often used to model the number of events occurring within a fixed period of time.
- ▶ **Model Assumption:** Events occur independently at a constant rate.
- ▶ Common applications: Counting the number of occurrences of events (e.g., arrival of customers, mutation occurrences in a DNA sequence).

Bayesian Poisson Model

- ▶ In Bayesian inference, we combine prior beliefs about a parameter with the likelihood of observed data.
- ▶ **Likelihood:** Assuming y events observed follows $y \sim \text{Poisson}(\lambda)$.
- ▶ **Prior for λ :** Typically a Gamma distribution due to its conjugacy, $\lambda \sim \text{Gamma}(a, b)$.
 - ▶ A non-informative prior can be implemented by setting a and b very close to zero (e.g., 0.001). This results in a very flat distribution, indicating high uncertainty and allowing the data to have a stronger influence on the posterior.
 - ▶ The flatness ensures that the prior does not overly constrain the likelihood, providing minimal initial information about the parameter.

Poisson Model in RStan

Stan Code (poisson_model.stan):

```
data {  
  int<lower=0> N;           // Number of events observed  
  int y[N];                // Observed counts  
}  
parameters {  
  real<lower=0> lambda;    // Rate parameter of Poisson  
}  
model {  
  // Non-informative prior: lambda ~ gamma(0.001, 0.001)  
  // Informative prior: lambda ~ gamma(9, 0.5)  
  lambda ~ gamma(9, 0.5);  // Change as needed  
  y ~ poisson(lambda);  
}
```


Poisson Model in RStan

R Code to Run Stan:

```
library(rstan)
data_obs <- c(3, 2, 1, 5, 4) // Example data
stan_data <- list(N = length(data_obs), y = data_obs)
fit <- stan(file = "poisson_model.stan", data = stan_data,
            iter = 2000, chains = 4)
print(fit) // Summary of posterior
plot(fit) // Visualize posterior distribution
```