

Pandas Library

Dhafer Malouche

February 16, 2022

<https://dhafermalouche.net>

Contents

- 1 Description from the Pandas documentation:
- 2 Series and DataFrames
 - 2.1 The Panda Series
 - 2.2 Pandas DataFrames
- 3 Data Manipulation

1 Description from the Pandas documentation:

- Pandas is a data analysis library providing fast, flexible, and expressive data structures designed to work with relational or table-like data (SQL table or Excel spreadsheet). It is a fundamental high-level building block for doing practical, real world data analysis in Python.
- Pandas is well suited for:
 - Tabular data with heterogeneously-typed columns, as in an SQL table or Excel spreadsheet Ordered and unordered (not necessarily fixed-frequency) time series data.
 - Arbitrary matrix data (homogeneously typed or heterogeneous) with row and column labels
 - Any other form of observational / statistical data sets.
- The data used with Pandas actually doesn't need be labeled at all to be placed into a Pandas data structure.
- The two primary data structures of Pandas, Series (1-dimensional) and DataFrame (2-dimensional), handle the vast majority of typical use cases in finance, statistics, social science, and many areas of engineering.
- Pandas is built **on top of** NumPy and is intended to integrate well within a scientific computing environment with many other 3rd party libraries.

Here are just a few of the things that Pandas does well:

- Easy handling of missing data (represented as NaN) in floating point as well as non-floating point data
- Size mutability: columns can be inserted and deleted from DataFrame and higher dimensional objects

- Automatic and explicit data alignment: objects can be explicitly aligned to a set of labels, or the user can simply ignore the labels and let Series, DataFrame, etc. automatically align the data for you in computations
- Powerful, flexible group by functionality to perform split-apply-combine operations on data sets, for both aggregating and transforming data
- Make it easy to convert ragged, differently-indexed data in other Python and NumPy data structures into DataFrame objects
- Intelligent label-based slicing, fancy indexing, and subsetting of large data sets
- Intuitive merging and joining data sets
- Flexible reshaping and pivoting of data sets
- Hierarchical labeling of axes (possible to have multiple labels per tick)
- Robust IO tools for loading data from flat files (CSV and delimited), Excel files, databases, and saving / loading data from the ultrafast [HDF5 format](#)
- Time series-specific functionality: date range generation and frequency conversion, moving window statistics, moving window linear regressions, date shifting and lagging, etc.

2 Series and DataFrames

We should first import Pandas into Python after installing it from the CMD prompt: `pip install pandas`

```
[1]: import pandas as pd
```

3 The Panda Series

The Series data structure in Pandas is a one-dimensional labeled array. + Data in the array can be of any type (integers, strings, floating point numbers, Python objects, etc.). + Data within the array is **homogeneous** + Pandas Series objects always have an index: this gives them both ndarray-like and dict-like properties.

Creating a Panda Serie: + Creation from a list + Creation from a dictionary + Creation from a ndarray + From an external source like a file (.csv,xls...)

From a list

```
[2]: temperature = [34, 56, 15, -9, -121, -5, 39]
days = ['Mon', 'Tue', 'Wed', 'Thu', 'Fri', 'Sat', 'Sun']

# create series
series_from_list = pd.Series(temperature, index=days)
series_from_list
```

```
[2]: Mon      34
     Tue      56
     Wed      15
     Thu      -9
     Fri     -121
     Sat      -5
```

```
Sun      39
dtype: int64
```

The series should contains homogeneous types

```
[3]: temperature = [34, 56, 'a', -9, -121, -5, 39]
     days = ['Mon', 'Tue', 'Wed', 'Thu', 'Fri', 'Sat', 'Sun']
```

We create series

```
[4]: series_from_list = pd.Series(temperature, index=days)
     series_from_list
```

```
[4]: Mon      34
     Tue      56
     Wed       a
     Thu      -9
     Fri     -121
     Sat      -5
     Sun      39
     dtype: object
```

from a dictionary

```
[5]: my_dict = {'Mon': 33, 'Tue': 19, 'Wed': 15, 'Thu': 89, 'Fri': 11, 'Sat': -5,
     ↪ 'Sun': 9}
     my_dict
```

```
[5]: {'Mon': 33, 'Tue': 19, 'Wed': 15, 'Thu': 89, 'Fri': 11, 'Sat': -5, 'Sun': 9}
```

```
[6]: series_from_dict = pd.Series(my_dict)
     series_from_dict
```

```
[6]: Mon      33
     Tue      19
     Wed      15
     Thu      89
     Fri      11
     Sat      -5
     Sun       9
     dtype: int64
```

From a numpy array

```
[7]: import numpy as np
```

I'm using linspace to create an array with spaced numbers over a specified interval: 15 numbers between 0 and 10

```
[8]: my_array = np.linspace(0,10,15)
      my_array
```

```
[8]: array([ 0.          ,  0.71428571,  1.42857143,  2.14285714,  2.85714286,
           3.57142857,  4.28571429,  5.          ,  5.71428571,  6.42857143,
           7.14285714,  7.85714286,  8.57142857,  9.28571429, 10.          ])
```

```
[9]: len(my_array)
```

```
[9]: 15
```

The array **must** be with dimension 1

```
[10]: series_from_ndarray = pd.Series(my_array)
      series_from_ndarray
```

```
[10]: 0      0.000000
      1      0.714286
      2      1.428571
      3      2.142857
      4      2.857143
      5      3.571429
      6      4.285714
      7      5.000000
      8      5.714286
      9      6.428571
     10      7.142857
     11      7.857143
     12      8.571429
     13      9.285714
     14     10.000000
      dtype: float64
```

4 Pandas DataFrames

DataFrame is a 2-dimensional labeled data structure with **columns** of potentially different types. You can think of it like a spreadsheet or SQL table, or a dict of Series objects. You can create a DataFrame from: + Dict of 1D ndarrays, lists, dicts, or Series + 2-D numpy.ndarray + From text, CSV, Excel files or databases + Many other ways

Reading the data.

Sample data: HR Employee Attrition and Performance You can get it from here and add it to your working directory:

<https://www.ibm.com/communities/analytics/watson-analytics-blog/hr-employee-attrition/>

Importing the xlsx file by considering the variable EmployeeNumber as an Index variable

```
[11]: # If Kaggle use this after uploading the xlsx into Kaggle
      ## data = pd.read_excel(io=" ../input/WA_Fn-UseC_-HR-Employee-Attrition.xlsx",
      ↪sheetname=0, index_col='EmployeeNumber')
```

```
pd.read_excel(io = "path to your excel data file, index_col = 'name of the column containing the row numbers or indexes')
```

Types of the variables

```
[12]: data.dtypes
```

```
[12]: Age                int64
      Attrition          object
      BusinessTravel     object
      DailyRate          int64
      Department         object
      DistanceFromHome   int64
      Education          int64
      EducationField     object
      EmployeeCount      int64
      EnvironmentSatisfaction int64
      Gender             object
      HourlyRate         int64
      JobInvolvement     int64
      JobLevel           int64
      JobRole            object
      JobSatisfaction    int64
      MaritalStatus      object
      MonthlyIncome      int64
      MonthlyRate        int64
      NumCompaniesWorked int64
      Over18             object
      OverTime           object
      PercentSalaryHike  int64
      PerformanceRating  int64
      RelationshipSatisfaction int64
      StandardHours      int64
      StockOptionLevel   int64
      TotalWorkingYears  int64
      TrainingTimesLastYear int64
      WorkLifeBalance     int64
      YearsAtCompany     int64
      YearsInCurrentRole  int64
      YearsSinceLastPromotion int64
      YearsWithCurrManager int64
      dtype: object
```

A preview of the data (the first 3 rows)

```
[39]: data.head(3)
```

```
[39]:
```

EmployeeNumber	Age	Attrition	BusinessTravel	DailyRate	\
1	41	Yes	Travel_Rarely	1102	
2	49	No	Travel_Frequently	279	
4	37	Yes	Travel_Rarely	1373	

EmployeeNumber	Department	DistanceFromHome	Education	\
1	Sales	1	2	
2	Research & Development	8	1	
4	Research & Development	2	2	

EmployeeNumber	EducationField	EmployeeCount	EnvironmentSatisfaction	...	\
1	Life Sciences	1	2	...	
2	Life Sciences	1	3	...	
4	Other	1	4	...	

EmployeeNumber	RelationshipSatisfaction	StandardHours	StockOptionLevel	\
1	1	80	0	
2	4	80	1	
4	2	80	0	

EmployeeNumber	TotalWorkingYears	TrainingTimesLastYear	WorkLifeBalance	\
1	8	0	1	
2	10	3	3	
4	7	3	3	

EmployeeNumber	YearsAtCompany	YearsInCurrentRole	YearsSinceLastPromotion	\
1	6	4	0	
2	10	7	1	
4	0	0	0	

EmployeeNumber	YearsWithCurrManager
1	5
2	7
4	0

```
[3 rows x 34 columns]
```

Name of the columns in the imported data.

```
[40]: data.columns
```

```
[40]: Index(['Age', 'Attrition', 'BusinessTravel', 'DailyRate', 'Department',  
        'DistanceFromHome', 'Education', 'EducationField', 'EmployeeCount',  
        'EnvironmentSatisfaction', 'Gender', 'HourlyRate', 'JobInvolvement',  
        'JobLevel', 'JobRole', 'JobSatisfaction', 'MaritalStatus',  
        'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked', 'Over18',  
        'OverTime', 'PercentSalaryHike', 'PerformanceRating',  
        'RelationshipSatisfaction', 'StandardHours', 'StockOptionLevel',  
        'TotalWorkingYears', 'TrainingTimesLastYear', 'WorkLifeBalance',  
        'YearsAtCompany', 'YearsInCurrentRole', 'YearsSinceLastPromotion',  
        'YearsWithCurrManager'],  
        dtype='object')
```

The preview of the variable Attrition

```
[44]: data['Attrition'].head()
```

```
[44]: EmployeeNumber  
1      Yes  
2      No  
4      Yes  
5      No  
7      No  
Name: Attrition, dtype: object
```

5 Data Manipulation

Selecting some variables from the original data and displaying a preview.

```
[45]: data[['Age', 'Gender', 'YearsAtCompany']].head()
```

```
[45]:
```

	Age	Gender	YearsAtCompany
EmployeeNumber			
1	41	Female	6
2	49	Male	10
4	37	Male	0
5	33	Female	8
7	27	Male	2

Creating a new variables. Transforming the Age in years to the Age in months.

```
[46]: data['AgeInMonths'] = 12*data['Age']  
data['AgeInMonths'].head()
```

```
[46]: EmployeeNumber  
1      492  
2      588  
4      444  
5      396  
7      324
```

Name: AgeInMonths, dtype: int64

Deleting the new created variable

```
[47]: del data['AgeInMonths']
```

```
[48]: data.columns
```

```
[48]: Index(['Age', 'Attrition', 'BusinessTravel', 'DailyRate', 'Department',  
        'DistanceFromHome', 'Education', 'EducationField', 'EmployeeCount',  
        'EnvironmentSatisfaction', 'Gender', 'HourlyRate', 'JobInvolvement',  
        'JobLevel', 'JobRole', 'JobSatisfaction', 'MaritalStatus',  
        'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked', 'Over18',  
        'OverTime', 'PercentSalaryHike', 'PerformanceRating',  
        'RelationshipSatisfaction', 'StandardHours', 'StockOptionLevel',  
        'TotalWorkingYears', 'TrainingTimesLastYear', 'WorkLifeBalance',  
        'YearsAtCompany', 'YearsInCurrentRole', 'YearsSinceLastPromotion',  
        'YearsWithCurrManager'],  
        dtype='object')
```

Extracting the some observations from on specific variable

```
[50]: data['BusinessTravel'][10:15]
```

```
[50]: EmployeeNumber  
14    Travel_Rarely  
15    Travel_Rarely  
16    Travel_Rarely  
18    Travel_Rarely  
19    Travel_Rarely  
Name: BusinessTravel, dtype: object
```

Extracting some rows from the whole dataframe

```
[52]: data[10:15]
```

```
[52]:
```

	Age	Attrition	BusinessTravel	DailyRate	\
EmployeeNumber					
14	35	No	Travel_Rarely	809	
15	29	No	Travel_Rarely	153	
16	31	No	Travel_Rarely	670	
18	34	No	Travel_Rarely	1346	
19	28	Yes	Travel_Rarely	103	

	Department	DistanceFromHome	Education	\
EmployeeNumber				
14	Research & Development	16	3	
15	Research & Development	15	2	
16	Research & Development	26	1	
18	Research & Development	19	2	
19	Research & Development	24	3	

EmployeeNumber	EducationField	EmployeeCount	EnvironmentSatisfaction	...	\
14	Medical	1	1	...	
15	Life Sciences	1	4	...	
16	Life Sciences	1	1	...	
18	Medical	1	2	...	
19	Life Sciences	1	3	...	

EmployeeNumber	RelationshipSatisfaction	StandardHours	StockOptionLevel	\
14	3	80	1	
15	4	80	0	
16	4	80	1	
18	3	80	1	
19	2	80	0	

EmployeeNumber	TotalWorkingYears	TrainingTimesLastYear	WorkLifeBalance	\
14	6	5	3	
15	10	3	3	
16	5	1	2	
18	3	2	3	
19	6	4	3	

EmployeeNumber	YearsAtCompany	YearsInCurrentRole	YearsSinceLastPromotion	\
14	5	4	0	
15	9	5	0	
16	5	2	4	
18	2	2	1	
19	4	2	0	

EmployeeNumber	YearsWithCurrManager
14	3
15	8
16	3
18	2
19	3

[5 rows x 34 columns]

Selecting specific rows from the index variable EmployeeNumbers

```
[57]: selected_EmployeeNumbers = [15, 94, 337, 1120]
```

```
[58]: data['YearsAtCompany']
```

```
[58]: EmployeeNumber
```

```
1      6
2     10
4       0
5       8
7       2
```

```
..
2061    5
2062    7
2064    6
2065    9
2068    4
```

```
Name: YearsAtCompany, Length: 1470, dtype: int64
```

```
[59]: data['YearsAtCompany'].loc[selected_EmployeeNumbers]
```

```
[59]: EmployeeNumber
```

```
15      9
94      5
337     2
1120    7
```

```
Name: YearsAtCompany, dtype: int64
```

```
[60]: data.loc[selected_EmployeeNumbers]
```

```
[60]:      Age Attrition      BusinessTravel      DailyRate \
```

```
EmployeeNumber
```

```
15      29      No      Travel_Rarely      153
94      29      No      Travel_Rarely     1328
337     31      No  Travel_Frequently     1327
1120    29      No      Travel_Rarely     1107
```

```
      Department      DistanceFromHome      Education \
```

```
EmployeeNumber
```

```
15      Research & Development      15      2
94      Research & Development      2      3
337     Research & Development      3      4
1120    Research & Development     28      4
```

```
      EducationField      EmployeeCount      EnvironmentSatisfaction ... \
```

```
EmployeeNumber
```

```
15      Life Sciences      1      4 ...
94      Life Sciences      1      3 ...
337      Medical      1      2 ...
1120    Life Sciences      1      3 ...
```

```
      RelationshipSatisfaction      StandardHours      StockOptionLevel \
```

```
EmployeeNumber
```

```
15      4      80      0
```

94	4	80	1
337	1	80	1
1120	1	80	1

EmployeeNumber	TotalWorkingYears	TrainingTimesLastYear	WorkLifeBalance	\
15	10	3	3	
94	6	3	3	
337	9	3	3	
1120	11	1	3	

EmployeeNumber	YearsAtCompany	YearsInCurrentRole	YearsSinceLastPromotion	\
15	9	5	0	
94	5	4	0	
337	2	2	2	
1120	7	5	1	

EmployeeNumber	YearsWithCurrManager
15	8
94	4
337	2
1120	7

[4 rows x 34 columns]

What's the YearsAtCompany of the row with EmployeeNumber equal to 94?

```
[62]: data.loc[94, 'YearsAtCompany']
```

```
[62]: 5
```

Frequency of the variable Department


```
[64]: data['Department'].value_counts()
```

```
[64]: Research & Development    961
Sales                          446
Human Resources                 63
Name: Department, dtype: int64
```

A barplot of the variable Department

```
[66]: data['Department'].value_counts().plot(kind='barh', title='Department')
```


```
[66]: <AxesSubplot:title={'center': 'Department'}>
```



Creating a pie chart

```
[67]: data['Department'].value_counts().plot(kind='pie', title='Department')
```

```
[67]: <AxesSubplot:title={'center': 'Department'}, ylabel='Department'>
```



Frequency of the variable Attrition

```
[70]: data['Attrition'].value_counts()
```

```
[70]: No      1233  
      Yes      237  
      Name: Attrition, dtype: int64
```

Frequency in percentage

```
[72]: data['Attrition'].value_counts(normalize=True)
```

```
[72]: No      0.838776  
      Yes      0.161224  
      Name: Attrition, dtype: float64
```

Compute the average of the variable HourlyRate

```
[73]: data['HourlyRate'].mean()
```

```
[73]: 65.89115646258503
```

What's the overall satisfaction of the Employees?

```
[75]: data['JobSatisfaction'].head()
```

```
[75]: EmployeeNumber  
      1      4  
      2      2  
      4      3  
      5      3  
      7      2  
      Name: JobSatisfaction, dtype: int64
```

Let us change the levels of the variable satisfaction by creating first a dictionary

```
[77]: JobSatisfaction_cat = {  
      1: 'Low',  
      2: 'Medium',  
      3: 'High',  
      4: 'Very High'  
      }
```

```
[78]: data['JobSatisfaction'] = data['JobSatisfaction'].map(JobSatisfaction_cat)  
      data['JobSatisfaction'].head()
```

```
[78]: EmployeeNumber  
      1      Very High  
      2      Medium  
      4      High  
      5      High  
      7      Medium  
      Name: JobSatisfaction, dtype: object
```

```
[79]: data['JobSatisfaction'].value_counts()
```

```
[79]: Very High    459
      High        442
      Low         289
      Medium      280
      Name: JobSatisfaction, dtype: int64
```


Computing percentages

```
[81]: 100*data['JobSatisfaction'].value_counts(normalize=True)
```

```
[81]: Very High    31.224490
      High        30.068027
      Low         19.659864
      Medium      19.047619
      Name: JobSatisfaction, dtype: float64
```

```
[82]: data['JobSatisfaction'].value_counts(normalize=True).plot(kind='pie',
      ↪title='Department')
```

```
[82]: <AxesSubplot:title={'center': 'Department'}, ylabel='JobSatisfaction'>
```



```
[88]: from pandas.api.types import CategoricalDtype
      cats=['Low', 'Medium', 'High', 'Very High']
      cat_type = CategoricalDtype(categories=cats, ordered=True)
      data['JobSatisfaction'] = data['JobSatisfaction'].astype(cat_type)
```

```
[89]: data['JobSatisfaction'].head()
```


```
[89]: EmployeeNumber
      1    Very High
```

```
2      Medium
4      High
5      High
7      Medium
Name: JobSatisfaction, dtype: category
Categories (4, object): ['Low' < 'Medium' < 'High' < 'Very High']
```

Sorting by frequencies (it's the default option) -

```
[91]: data['JobSatisfaction'].value_counts().plot(kind='barh', title='Department')
```


```
[91]: <AxesSubplot:title={'center':'Department'}>
```



Canceling the default sorting option and the bars will be sorted according to the categories

```
[92]: data['JobSatisfaction'].value_counts(sort=False).plot(kind='barh', title='Department')
```

```
[92]: <AxesSubplot:title={'center':'Department'}>
```



```
[93]: data['JobSatisfaction'] == 'Low'
```

```
[93]: EmployeeNumber
1      False
2      False
4      False
5      False
7      False
...
2061   False
2062    True
2064   False
2065   False
2068   False
Name: JobSatisfaction, Length: 1470, dtype: bool
```

```
[94]: data.loc[data['JobSatisfaction'] == 'Low'].index
```

```
[94]: Int64Index([ 10,  20,  27,  31,  33,  38,  51,  52,  54,  68,
...
1975, 1980, 1998, 2021, 2023, 2038, 2054, 2055, 2057, 2062],
              dtype='int64', name='EmployeeNumber', length=289)
```

```
[95]: data['JobInvolvement'].head()
```

```
[95]: EmployeeNumber
1      3
2      2
```



```
4    2
5    3
7    3
```

```
Name: JobInvolvement, dtype: int64
```

Selecting observation of a specific interest: Those with either “Low” or “Very High” Job satisfaction

```
[107]: subset_of_interest = data.loc[(data['JobSatisfaction'] == "Low") |
    ↳ (data['JobSatisfaction'] == "Very High")]
subset_of_interest.shape
```

```
[107]: (748, 34)
```

```
[108]: subset_of_interest.head()
```

```
[108]:
```

	Age	Attrition	BusinessTravel	DailyRate	\
EmployeeNumber					
1	41	Yes	Travel_Rarely	1102	
8	32	No	Travel_Frequently	1005	
10	59	No	Travel_Rarely	1324	
18	34	No	Travel_Rarely	1346	
20	29	No	Travel_Rarely	1389	

	Department	DistanceFromHome	Education	\
EmployeeNumber				
1	Sales		1	2
8	Research & Development		2	2
10	Research & Development		3	3
18	Research & Development		19	2
20	Research & Development		21	4

	EducationField	EmployeeCount	EnvironmentSatisfaction	...	\
EmployeeNumber				...	
1	Life Sciences	1		2	...
8	Life Sciences	1		4	...
10	Medical	1		3	...
18	Medical	1		2	...
20	Life Sciences	1		2	...

	RelationshipSatisfaction	StandardHours	StockOptionLevel	\
EmployeeNumber				
1	1	80	0	
8	3	80	0	
10	1	80	3	
18	3	80	1	
20	3	80	1	

	TotalWorkingYears	TrainingTimesLastYear	WorkLifeBalance	\
EmployeeNumber				

1	8	0	1
8	8	2	2
10	12	3	2
18	3	2	3
20	10	1	3

	YearsAtCompany	YearsInCurrentRole	YearsSinceLastPromotion	\
EmployeeNumber				
1	6	4		0
8	7	7		3
10	1	0		0
18	2	2		1
20	10	9		8

	YearsWithCurrManager
EmployeeNumber	
1	5
8	6
10	0
18	2
20	8

[5 rows x 34 columns]

```
[109]: subset_of_interest['JobSatisfaction'].value_counts()
```

```
[109]: Very High    459
Low            289
Medium         0
High           0
Name: JobSatisfaction, dtype: int64
```

Let's then remove the categories or levels that we won't use

```
[110]: subset_of_interest['JobSatisfaction'].cat.remove_unused_categories(inplace=True)
```

```
C:\ProgramData\Anaconda3\lib\site-
packages\pandas\core\arrays\categorical.py:2631: FutureWarning: The `inplace`
parameter in pandas.Categorical.remove_unused_categories is deprecated and will
be removed in a future version.
    res = method(*args, **kwargs)
```

The categories 'Medium' and 'High' won't be displayed

```
[112]: subset_of_interest['JobSatisfaction'].value_counts()
```

```
[112]: Very High    459
Low            289
Name: JobSatisfaction, dtype: int64
```

```
[113]: grouped = subset_of_interest.groupby('JobSatisfaction')
```

```
[116]: grouped.head()
```

```
[116]:
```

	Age	Attrition	BusinessTravel	DailyRate	\
EmployeeNumber					
1	41	Yes	Travel_Rarely	1102	
8	32	No	Travel_Frequently	1005	
10	59	No	Travel_Rarely	1324	
18	34	No	Travel_Rarely	1346	
20	29	No	Travel_Rarely	1389	
22	22	No	Non-Travel	1123	
23	53	No	Travel_Rarely	1219	
27	36	Yes	Travel_Rarely	1218	
31	34	Yes	Travel_Rarely	699	
33	32	Yes	Travel_Frequently	1125	

	Department	DistanceFromHome	Education	\
EmployeeNumber				
1	Sales	1	2	
8	Research & Development	2	2	
10	Research & Development	3	3	
18	Research & Development	19	2	
20	Research & Development	21	4	
22	Research & Development	16	2	
23	Sales	2	4	
27	Sales	9	4	
31	Research & Development	6	1	
33	Research & Development	16	1	

	EducationField	EmployeeCount	EnvironmentSatisfaction	...	\
EmployeeNumber				...	
1	Life Sciences	1	2	...	
8	Life Sciences	1	4	...	
10	Medical	1	3	...	
18	Medical	1	2	...	
20	Life Sciences	1	2	...	
22	Medical	1	4	...	
23	Life Sciences	1	1	...	
27	Life Sciences	1	3	...	
31	Medical	1	2	...	
33	Life Sciences	1	2	...	

	RelationshipSatisfaction	StandardHours	StockOptionLevel	\
EmployeeNumber				
1	1	80	0	
8	3	80	0	
10	1	80	3	
18	3	80	1	
20	3	80	1	

22	2	80	2
23	3	80	0
27	2	80	0
31	3	80	0
33	2	80	0

EmployeeNumber	TotalWorkingYears	TrainingTimesLastYear	WorkLifeBalance	\
1	8	0	1	
8	8	2	2	
10	12	3	2	
18	3	2	3	
20	10	1	3	
22	1	2	2	
23	31	3	3	
27	10	4	3	
31	8	2	3	
33	10	5	3	

EmployeeNumber	YearsAtCompany	YearsInCurrentRole	YearsSinceLastPromotion	\
1	6	4	0	
8	7	7	3	
10	1	0	0	
18	2	2	1	
20	10	9	8	
22	1	0	0	
23	25	8	3	
27	5	3	0	
31	4	2	1	
33	10	2	6	

EmployeeNumber	YearsWithCurrManager
1	5
8	6
10	0
18	2
20	8
22	0
23	7
27	3
31	3
33	7

[10 rows x 34 columns]

```
[114]: grouped.groups
```

```
[114]: {'Low': [10, 20, 27, 31, 33, 38, 51, 52, 54, 68, 70, 74, 75, 81, 86, 88, 100,
101, 113, 124, 133, 134, 145, 153, 170, 190, 197, 199, 200, 235, 239, 240, 241,
244, 250, 267, 274, 282, 288, 297, 299, 303, 328, 334, 339, 340, 347, 351, 362,
369, 374, 382, 390, 396, 412, 424, 425, 429, 451, 454, 474, 486, 510, 515, 517,
522, 524, 530, 532, 534, 536, 538, 549, 567, 573, 590, 605, 615, 625, 630, 648,
650, 662, 664, 667, 682, 684, 702, 705, 725, 728, 729, 732, 733, 742, 758, 764,
771, 775, 776, ...], 'Very High': [1, 8, 18, 22, 23, 24, 30, 36, 39, 40, 42, 45,
49, 53, 57, 62, 63, 72, 73, 76, 78, 79, 97, 98, 104, 106, 107, 112, 116, 117,
118, 120, 137, 139, 140, 143, 144, 148, 152, 154, 155, 158, 165, 169, 174, 179,
184, 192, 195, 198, 207, 215, 217, 221, 223, 228, 230, 242, 243, 245, 246, 262,
264, 273, 275, 281, 283, 286, 287, 291, 298, 302, 306, 309, 311, 312, 315, 316,
319, 323, 325, 327, 333, 335, 336, 338, 346, 349, 353, 361, 367, 372, 373, 377,
378, 380, 388, 389, 391, 393, ...]}
```

The Low satisfaction group

```
[115]: grouped.get_group('Low').head()
```

```
[115]:
```

	Age	Attrition	BusinessTravel	DailyRate	\
EmployeeNumber					
10	59	No	Travel_Rarely	1324	
20	29	No	Travel_Rarely	1389	
27	36	Yes	Travel_Rarely	1218	
31	34	Yes	Travel_Rarely	699	
33	32	Yes	Travel_Frequently	1125	

	Department	DistanceFromHome	Education	\
EmployeeNumber				
10	Research & Development	3	3	
20	Research & Development	21	4	
27	Sales	9	4	
31	Research & Development	6	1	
33	Research & Development	16	1	

	EducationField	EmployeeCount	EnvironmentSatisfaction	...	\
EmployeeNumber				...	
10	Medical	1	3	...	
20	Life Sciences	1	2	...	
27	Life Sciences	1	3	...	
31	Medical	1	2	...	
33	Life Sciences	1	2	...	

	RelationshipSatisfaction	StandardHours	StockOptionLevel	\
EmployeeNumber				
10	1	80	3	
20	3	80	1	
27	2	80	0	

31	3	80	0
33	2	80	0

	TotalWorkingYears	TrainingTimesLastYear	WorkLifeBalance	\
EmployeeNumber				
10	12	3	2	
20	10	1	3	
27	10	4	3	
31	8	2	3	
33	10	5	3	

	YearsAtCompany	YearsInCurrentRole	YearsSinceLastPromotion	\
EmployeeNumber				
10	1	0	0	
20	10	9	8	
27	5	3	0	
31	4	2	1	
33	10	2	6	

	YearsWithCurrManager
EmployeeNumber	
10	0
20	8
27	3
31	3
33	7

[5 rows x 34 columns]

and the Very High satisfaction group

```
[104]: grouped.get_group('Very High').head()
```

	Age	Attrition	BusinessTravel	DailyRate	\
EmployeeNumber					
1	41	Yes	Travel_Rarely	1102	
8	32	No	Travel_Frequently	1005	
18	34	No	Travel_Rarely	1346	
22	22	No	Non-Travel	1123	
23	53	No	Travel_Rarely	1219	

	Department	DistanceFromHome	Education	\
EmployeeNumber				
1	Sales	1	2	
8	Research & Development	2	2	
18	Research & Development	19	2	
22	Research & Development	16	2	
23	Sales	2	4	

EmployeeNumber	EducationField	EmployeeCount	EnvironmentSatisfaction	...	\
1	Life Sciences	1	2	...	
8	Life Sciences	1	4	...	
18	Medical	1	2	...	
22	Medical	1	4	...	
23	Life Sciences	1	1	...	

EmployeeNumber	RelationshipSatisfaction	StandardHours	StockOptionLevel	\
1	1	80	0	
8	3	80	0	
18	3	80	1	
22	2	80	2	
23	3	80	0	

EmployeeNumber	TotalWorkingYears	TrainingTimesLastYear	WorkLifeBalance	\
1	8	0	1	
8	8	2	2	
18	3	2	3	
22	1	2	2	
23	31	3	3	

EmployeeNumber	YearsAtCompany	YearsInCurrentRole	YearsSinceLastPromotion	\
1	6	4	0	
8	7	7	3	
18	2	2	1	
22	1	0	0	
23	25	8	3	

EmployeeNumber	YearsWithCurrManager
1	5
8	6
18	2
22	0
23	7

[5 rows x 34 columns]

The average of the Age of each group

```
[120]: grouped[['Age', 'JobSatisfaction']].head()
```

```
[120]:
```

	Age	JobSatisfaction
EmployeeNumber		
1	41	Very High
8	32	Very High
10	59	Low
18	34	Very High
20	29	Low
22	22	Very High
23	53	Very High
27	36	Low
31	34	Low
33	32	Low

```
[121]: grouped['Age'].mean()
```

```
[121]: JobSatisfaction
Low      36.916955
Very High 36.795207
Name: Age, dtype: float64
```

```
[122]: grouped['Age'].describe()
```

```
[122]:
```


	count	mean	std	min	25%	50%	75%	max
JobSatisfaction								
Low	289.0	36.916955	9.245496	19.0	30.0	36.0	42.0	60.0
Very High	459.0	36.795207	9.125609	18.0	30.0	35.0	43.0	60.0

```
[ ]: grouped['Age'].describe().unstack()
```

Comparing densities

```
[124]: grouped['Age'].plot(kind='density', title='Age')
```

```
[124]: JobSatisfaction
Low      AxesSubplot(0.125,0.125;0.775x0.755)
Very High AxesSubplot(0.125,0.125;0.775x0.755)
Name: Age, dtype: object
```

By Department

```
[125]: grouped['Department'].value_counts().unstack()
```

```
[125]: Department      Human Resources  Research & Development  Sales
JobSatisfaction
Low                    11                192         86
Very High             17                295        147
```


We can normalize it

```
[126]: grouped['Department'].value_counts(normalize=True).unstack()
```

```
[126]: Department      Human Resources  Research & Development  Sales
JobSatisfaction
Low                    0.038062        0.664360  0.297578
Very High             0.037037        0.642702  0.320261
```


```
[127]: grouped['Department'].value_counts().unstack().plot(kind="barh")
```

```
[127]: <AxesSubplot:ylabel='JobSatisfaction'>
```



```
[128]: grouped['Department'].value_counts(normalize=True).unstack().plot(kind="barh")
```


```
[128]: <AxesSubplot:ylabel='JobSatisfaction'>
```



We can compare it with the whole sample


```
[129]: data['Department'].value_counts(normalize=True,sort=False).plot(kind="barh")
```

```
[129]: <AxesSubplot:>
```



```
[132]: data['Department'].value_counts(normalize=True,sort=False).plot(kind="barh")
```

```
[132]: <AxesSubplot:>
```



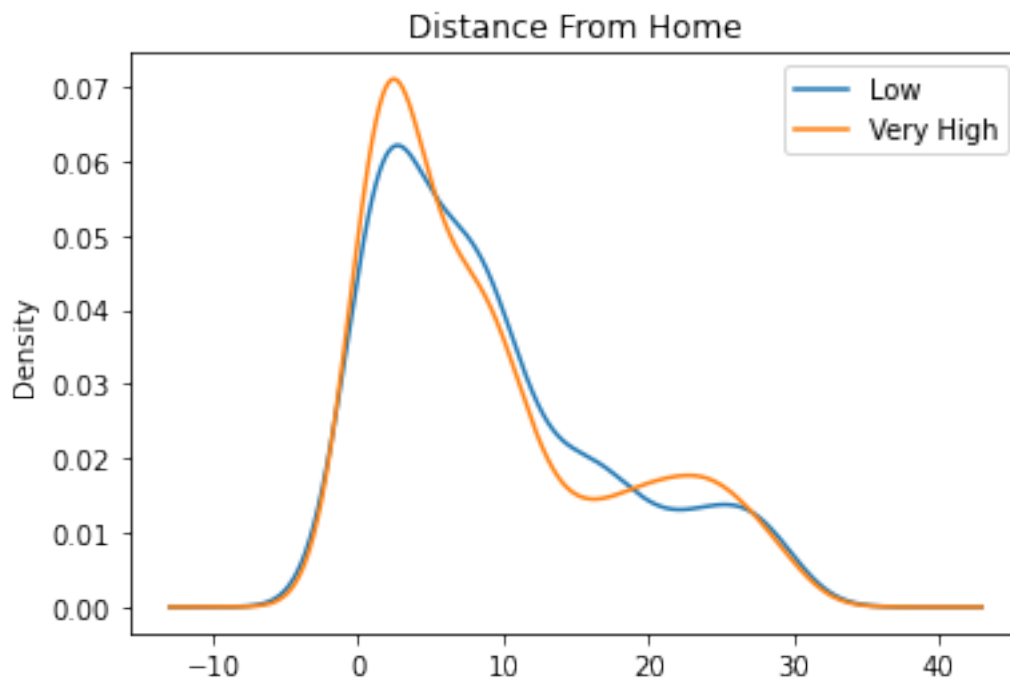
```
[133]: grouped['DistanceFromHome'].describe().unstack()
```

```
[133]:      JobSatisfaction  
count  Low      289.000000
```

	Very High	459.000000
mean	Low	9.190311
	Very High	9.030501
std	Low	8.045127
	Very High	8.257004
min	Low	1.000000
	Very High	1.000000
25%	Low	2.000000
	Very High	2.000000
50%	Low	7.000000
	Very High	7.000000
75%	Low	14.000000
	Very High	14.000000
max	Low	29.000000
	Very High	29.000000
dtype: float64		

```
[134]: grouped['DistanceFromHome'].plot(kind='density', title='Distance From Home',
      ↪      legend=True)
```

```
[134]: JobSatisfaction
Low      AxesSubplot(0.125,0.125;0.775x0.755)
Very High AxesSubplot(0.125,0.125;0.775x0.755)
Name: DistanceFromHome, dtype: object
```




```
[135]: grouped['HourlyRate'].describe()
```

```
[135]:
```

	count	mean	std	min	25%	50%	75%	max
JobSatisfaction								
Low	289.0	68.636678	20.439515	30.0	52.0	72.0	86.0	100.0
Very High	459.0	64.681917	20.647571	30.0	47.0	64.0	82.5	100.0

```
[136]: grouped['HourlyRate'].plot(kind='density', title='Hourly Rate',legend=True)
```

```
[136]: JobSatisfaction
Low      AxesSubplot(0.125,0.125;0.775x0.755)
Very High AxesSubplot(0.125,0.125;0.775x0.755)
Name: HourlyRate, dtype: object
```



```
[137]: grouped['MonthlyIncome'].describe()
```

```
[137]:
```


	count	mean	std	min	25%	50% \
JobSatisfaction						
Low	289.0	6561.570934	4645.170134	1091.0	3072.0	4968.0
Very High	459.0	6472.732026	4573.906428	1051.0	2927.5	5126.0

	75%	max
JobSatisfaction		
Low	8564.0	19943.0
Very High	7908.0	19845.0

```
[138]: grouped['HourlyRate'].plot(kind='density', title='Hourly Rate',legend=True)
```

```
[138]: JobSatisfaction
Low      AxesSubplot(0.125,0.125;0.775x0.755)
```

```
Very High    AxesSubplot(0.125,0.125;0.775x0.755)  
Name: HourlyRate, dtype: object
```



```
[13]: !pip install numpy
```

```
Defaulting to user installation because normal site-packages is not writeable  
Requirement already satisfied: numpy in  
c:\users\dhafe\appdata\roaming\python\python310\site-packages (1.22.1)
```

```
[ ]:
```

1